

Importing Libraries

```
In [1]: import pyforest #for importing all libraries at once
import pygwalker as pyg
import pycaret
```

```
In [2]: from sklearn.model_selection import cross_val_score, GridSearchCV, RandomizedSearchCV
from sklearn.tree import DecisionTreeClassifier
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from xgboost import XGBClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.neural_network import MLPClassifier
from sklearn.ensemble import GradientBoostingClassifier
from lightgbm import LGBMClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
#to ignore all the warnings
import warnings
warnings.filterwarnings("ignore")
```

Loading the Dataset and play with the data

```
In [3]: # Load the dataset
df1 = pd.read_csv("C&T train dataset.csv")
df2 = pd.read_csv("C&T test dataset.csv")
copy = pd.read_csv('C&T test dataset.csv')
#keep a copy of the test beacuse we will be later dropping the sno column but it wi
```

```
In [4]: #We can get the whole info about our data plus we can perform some visualisations
gwalker = pyg.walk(df1)
```

```
Box(children=(HTML(value='<div id="ifr-pyg-000616d423f52488k0v3qTUAo4wuSg6V" style="height: auto">\n    <head>...
```

```
In [5]: df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 800 entries, 0 to 799
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   sno                    800 non-null    int64
1   acc_info               800 non-null    object
2   duration_month         800 non-null    int64
3   credit_history         800 non-null    object
4   purpose                800 non-null    object
5   savings_acc            800 non-null    object
6   employment_st          792 non-null    object
7   poi                    788 non-null    float64
8   personal_status        800 non-null    object
9   gurantors              792 non-null    object
10  resident_since         800 non-null    int64
11  property_type          800 non-null    object
12  age                    796 non-null    float64
13  installment_type       800 non-null    object
14  housing_type           793 non-null    object
15  credits_no             800 non-null    int64
16  job_type               800 non-null    object
17  liables                800 non-null    int64
18  telephone              800 non-null    object
19  foreigner              800 non-null    object
20  Group_no               800 non-null    int64
dtypes: float64(2), int64(6), object(13)
memory usage: 131.4+ KB
```

EDA and Preprocessing

Handling null values and duplicates

Train dataset

```
In [6]: df1.duplicated().sum()    #no duplicate values
```

```
Out[6]: 0
```

```
In [7]: # Display the number of missing values in each column
print("Number of missing values in each column:")
print(df1.isnull().sum())

# Numeric Columns: Impute missing values with mean
numeric_columns = ['poi', 'age']
for col in numeric_columns:
    df1[col].fillna(df1[col].mean(), inplace=True)

# Categorical Columns: Impute missing values with mode
categorical_columns = ['employment_st', 'housing_type', 'gurantors']
for col in categorical_columns:
```

```

df1[col].fillna(df1[col].mode()[0], inplace=True)

# Display the number of missing values after handling
print("\nNumber of missing values after handling:")
print(df1.isnull().sum())

```

Number of missing values in each column:

```

sno                0
acc_info           0
duration_month     0
credit_history     0
purpose            0
savings_acc        0
employment_st      8
poi               12
personal_status    0
gurantors          8
resident_since     0
property_type      0
age                4
installment_type   0
housing_type       7
credits_no         0
job_type           0
liables            0
telephone          0
foreigner          0
Group_no           0
dtype: int64

```

Number of missing values after handling:

```

sno                0
acc_info           0
duration_month     0
credit_history     0
purpose            0
savings_acc        0
employment_st      0
poi               0
personal_status    0
gurantors          0
resident_since     0
property_type      0
age                0
installment_type   0
housing_type       0
credits_no         0
job_type           0
liables            0
telephone          0
foreigner          0
Group_no           0
dtype: int64

```

Test Dataset

```
In [8]: df2.duplicated().sum()    #no duplicate values
```

```
Out[8]: 0
```

```
In [9]: # Display the number of missing values in each column
print("Number of missing values in each column:")
print(df2.isnull().sum())

# Numeric Columns: Impute missing values with mean
numeric_columns = ['poi', 'age']
for col in numeric_columns:
    df2[col].fillna(df2[col].mean(), inplace=True)

# Categorical Columns: Impute missing values with mode
categorical_columns = ['employment_st', 'housing_type', 'gurantors']
for col in categorical_columns:
    df2[col].fillna(df2[col].mode()[0], inplace=True)

# Display the number of missing values after handling
print("\nNumber of missing values after handling:")
print(df2.isnull().sum())
```

Number of missing values in each column:

sno	0
acc_info	0
duration_month	0
credit_history	0
purpose	0
savings_acc	0
employment_st	5
poi	5
personal_status	0
gurantors	0
resident_since	0
property_type	0
age	2
installment_type	0
housing_type	8
credits_no	0
job_type	0
liables	0
telephone	0
foreigner	0
dtype: int64	

Number of missing values after handling:

sno	0
acc_info	0
duration_month	0
credit_history	0
purpose	0
savings_acc	0
employment_st	0
poi	0
personal_status	0
gurantors	0
resident_since	0
property_type	0
age	0
installment_type	0
housing_type	0
credits_no	0
job_type	0
liables	0
telephone	0
foreigner	0
dtype: int64	

Data Visualisation

```
In [10]: #Deleting the columns which are not necessary for our classification problem  
#df1=df1.drop(['sno','telephone','personal_status'],axis=1)  
#df2=df2.drop(['sno','telephone','personal_status'],axis=1)
```

```
In [11]: import matplotlib.pyplot as plt  
import seaborn as sns
```

```
plt.figure(figsize=(15,10))

plt.subplot(2,3,1)
sns.countplot(x='acc_info', hue='Group_no', data=df1, palette='plasma')
plt.xlabel('Account Information', fontsize=14)
plt.ylabel('Number', fontsize=14)

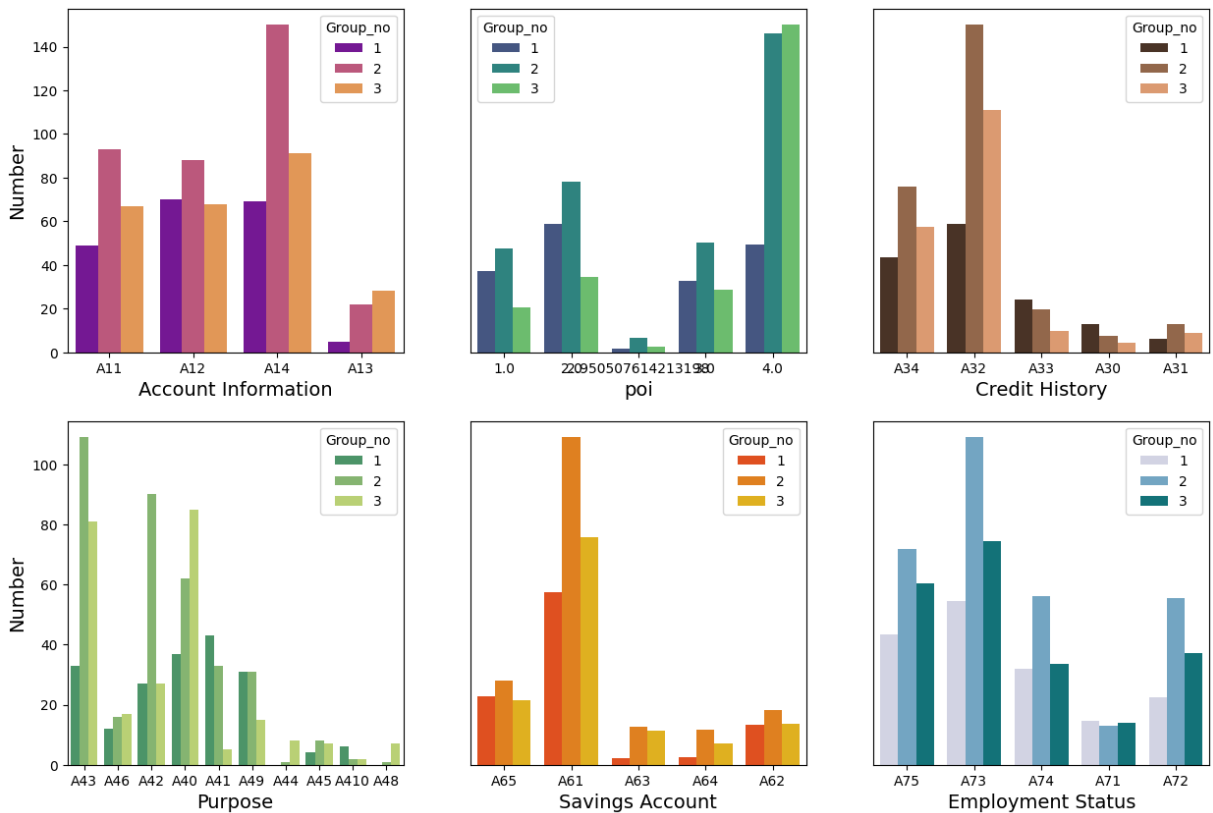
plt.subplot(2,3,2)
sns.countplot(x='poi', hue='Group_no', data=df1, palette='viridis')
plt.xlabel('poi', fontsize=14)
plt.ylabel(' ')
plt.yticks([])

plt.subplot(2,3,3)
sns.countplot(x='credit_history', hue='Group_no', data=df1, palette='copper')
plt.xlabel('Credit History', fontsize=14)
plt.ylabel(' ')
plt.yticks([])

plt.subplot(2,3,4)
sns.countplot(x='purpose', hue='Group_no', data=df1, palette='summer')
plt.xlabel('Purpose', fontsize=14)
plt.ylabel('Number', fontsize=14)

plt.subplot(2,3,5)
sns.countplot(x='savings_acc', hue='Group_no', data=df1, palette='autumn')
plt.xlabel('Savings Account', fontsize=14)
plt.ylabel(' ')
plt.yticks([])

plt.subplot(2,3,6)
sns.countplot(x='employment_st', hue='Group_no', data=df1, palette='PuBuGn')
plt.xlabel('Employment Status', fontsize=14)
plt.ylabel(' ')
plt.yticks([])
plt.show()
```



Univariate Analysis

```
In [12]: # Define categorical and numerical features
categorical_features = ['acc_info', 'credit_history', 'purpose', 'savings_acc', 'em
                    'gurantors', 'property_type', 'installment_type',
                    'housing_type', 'job_type', 'foreigner']
numerical_features = ['duration_month', 'poi', 'resident_since', 'age', 'credits_no
```

```
In [13]: # Select only numeric columns
df1_numeric = df1.select_dtypes(include=['number'])
df1_categorical = df1.select_dtypes(exclude=['number'])


# Now calculate correlations
#correlation_matrix = df1_numeric.corr()
```

```
In [14]: from scipy.stats import skew, kurtosis
key = []
skewval = []
kurtval = []
for i in df1_numeric :
    print('The skewness of the ', i, '=', skew(df1_numeric [i]))
    print('The kurtosis of the ', i, '=', kurtosis(df1_numeric [i]))
    plt.figure()
    sns.distplot(df1_numeric[i], color='blue', kde=True,
                hist_kws = {'color':'green', 'edgecolor':'black', 'linewidth':2, 'al
    plt.show()
    key.append(i)
```

```
skewval.append(skew(df1_numeric [i]))  
kurtval.append(kurtosis(df1_numeric [i]))
```


The skewness of the sno = 0.0

The kurtosis of the sno = -1.2000037500058593




The skewness of the duration_month = 1.1470027029543717


The kurtosis of the duration_month = 1.0529676149863754




The skewness of the poi = -0.5039066555236439
 The kurtosis of the poi = -1.2293262165867966




The skewness of the resident_since = -0.2767323024835812
 The kurtosis of the resident_since = -1.3766474938843547




The skewness of the age = 1.060425705742385
The kurtosis of the age = 0.6927026736563899




The skewness of the credits_no = 1.2652514728538597
The kurtosis of the credits_no = 1.4791338504629952



The skewness of the liabilities = 1.9881351927895627
 The kurtosis of the liabilities = 1.9526815448083923



The skewness of the Group_no = -0.12325434930243198
 The kurtosis of the Group_no = -1.1851006072150474



```
In [15]: dict = {'variables':key,'skew_values':skewval,'kurtosis_values':kurtval}
```

```
In [16]: show_val = pd.DataFrame(dict)
```

```
In [17]: show_val
```



```
Out[17]:
```



	variables	skew_values	kurtosis_values
0	sno	0.000000	-1.200004
1	duration_month	1.147003	1.052968
2	poi	-0.503907	-1.229326
3	resident_since	-0.276732	-1.376647
4	age	1.060426	0.692703
5	credits_no	1.265251	1.479134
6	liables	1.988135	1.952682
7	Group_no	-0.123254	-1.185101



Handling outliers



```
In [18]: for i in df1_numeric:
          plt.figure()
```

```
sns.boxplot(df1_numeric[i],color='blue')  
plt.show()
```









```
In [19]: #Impute the outliers from the data
# Calculate the first quartile (Q1) and third quartile (Q3)
Q1 = df1_numeric.quantile(0.25)
Q3 = df1_numeric.quantile(0.75)

# Calculate the IQR
```



```

IQR = Q3 - Q1

# Define the threshold for outliers (e.g., 1.5 times the IQR)
threshold = 1.5

# Impute outliers with the median
# The clip() method limits the values in the DataFrame to be within the specified limits
df1_imputed_median = df1_numeric.clip(lower=Q1 - threshold * IQR, upper=Q3 + threshold * IQR)

```

```

In [20]: # Combine imputed numeric columns with original categorical columns
df1_new = pd.concat([df1_imputed_median, df1_categorical], axis=1)

```

```

In [21]: df1_new

```

```

Out[21]:

```

	sno	duration_month	poi	resident_since	age	credits_no	liables	Group_no	acc_info
0	1	6	4.0	4	62.0	2.0	1	3	A1
1	2	42	2.0	2	22.0	1.0	1	1	A1
2	3	12	2.0	3	49.0	1.0	1	2	A1
3	4	42	2.0	4	45.0	1.0	1	1	A1
4	5	24	3.0	4	53.0	2.0	1	1	A1
...
795	796	9	2.0	4	22.0	1.0	1	2	A1
796	797	18	1.0	4	51.0	1.0	1	1	A1
797	798	12	2.0	4	22.0	2.0	1	3	A1
798	799	24	4.0	4	54.0	2.0	1	3	A1
799	800	9	4.0	2	35.0	1.0	1	2	A1

800 rows × 21 columns

Checking whether our dataset is imbalanced

```

In [22]: # Assuming 'Group_no' is the target variable
class_distribution = df1_new['Group_no'].value_counts()

print("Class Distribution:")
print(class_distribution)

# Calculate the imbalance ratio
imbalance_ratio = class_distribution.min() / class_distribution.max()

print("Imbalance Ratio:", imbalance_ratio)

```

```
Class Distribution:
Group_no
2    353
3    254
1    193
Name: count, dtype: int64
Imbalance Ratio: 0.546742209631728
```

So our dataset is Balanced

Bivariate Analysis

```
In [23]: # Now calculate correlations
correlation_matrix = df1_numeric.corr()
```


```
In [24]: correlation_matrix
```

```
Out[24]:
```

	sno	duration_month	poi	resident_since	age	credits_no
sno	1.000000	-0.001513	-0.000196	0.035694	-0.039311	0.001155
duration_month	-0.001513	1.000000	0.076066	0.030762	-0.045680	-0.013814
poi	-0.000196	0.076066	1.000000	0.039182	0.075055	0.014575
resident_since	0.035694	0.030762	0.039182	1.000000	0.251351	0.072087
age	-0.039311	-0.045680	0.075055	0.251351	1.000000	0.138340
credits_no	0.001155	-0.013814	0.014575	0.072087	0.138340	1.000000
liables	-0.004457	-0.023749	-0.066547	0.050136	0.129534	0.081980
Group_no	0.035358	-0.588725	0.259411	-0.026316	-0.003826	-0.056597

```
In [25]: sns.heatmap(correlation_matrix,annot=True,cmap='coolwarm')
```

```
Out[25]: <Axes: >
```



Label encoding

```
In [26]: # Create a Label encoder object
label_encoder = LabelEncoder()

# Iterate over columns and encode categorical columns
for column in df1_new.columns:
    if df1_new[column].dtype == 'object': # Check if column contains categorical data
        df1_new[column] = label_encoder.fit_transform(df1_new[column])

# Display the encoded DataFrame
df1_new
```

Out[26]:

	sno	duration_month	poi	resident_since	age	credits_no	liables	Group_no	acc_info
0	1	6	4.0	4	62.0	2.0	1	3	
1	2	42	2.0	2	22.0	1.0	1	1	
2	3	12	2.0	3	49.0	1.0	1	2	
3	4	42	2.0	4	45.0	1.0	1	1	
4	5	24	3.0	4	53.0	2.0	1	1	
...
795	796	9	2.0	4	22.0	1.0	1	2	
796	797	18	1.0	4	51.0	1.0	1	1	
797	798	12	2.0	4	22.0	2.0	1	3	
798	799	24	4.0	4	54.0	2.0	1	3	
799	800	9	4.0	2	35.0	1.0	1	2	

800 rows × 21 columns

In [27]: df2

Out[27]:

	sno	acc_info	duration_month	credit_history	purpose	savings_acc	employment_st
0	1	A14	24	A34	A46	A61	A75
1	2	A12	18	A34	A43	A61	A75
2	3	A11	20	A34	A42	A61	A75
3	4	A14	12	A34	A43	A65	A75
4	5	A12	12	A32	A40	A65	A71
...
195	196	A14	12	A32	A42	A61	A74
196	197	A11	30	A32	A41	A61	A73
197	198	A14	12	A32	A43	A61	A75
198	199	A11	45	A32	A43	A61	A73
199	200	A12	45	A34	A41	A62	A71

200 rows × 20 columns

In [28]:

```
# Create a Label encoder object
label_encoder = LabelEncoder()

# Iterate over columns and encode categorical columns
```

```

for column in df2.columns:
    if df2[column].dtype == 'object': # Check if column contains categorical data
        df2[column] = label_encoder.fit_transform(df2[column])

# Display the encoded DataFrame
df2

```

Out[28]:

	sno	acc_info	duration_month	credit_history	purpose	savings_acc	employment_st
0	1	3	24	4	7	0	4
1	2	1	18	4	4	0	4
2	3	0	20	4	3	0	4
3	4	3	12	4	4	4	4
4	5	1	12	2	0	4	0
...
195	196	3	12	2	3	0	3
196	197	0	30	2	1	0	2
197	198	3	12	2	4	0	4
198	199	0	45	2	4	0	2
199	200	1	45	4	1	1	0

200 rows × 20 columns

```

In [29]: # Get the number of categories in each column
category_counts = {}
for column in df1_new.columns:
    category_counts[column] = len(df1_new[column].value_counts())

# Print or use the counts as needed
for column, count in category_counts.items():
    print(f"Column '{column}' has {count} categories.")

```

Column 'sno' has 800 categories.
 Column 'duration_month' has 27 categories.
 Column 'poi' has 5 categories.
 Column 'resident_since' has 4 categories.
 Column 'age' has 45 categories.
 Column 'credits_no' has 4 categories.
 Column 'liables' has 1 categories.
 Column 'Group_no' has 3 categories.
 Column 'acc_info' has 4 categories.
 Column 'credit_history' has 5 categories.
 Column 'purpose' has 10 categories.
 Column 'savings_acc' has 5 categories.
 Column 'employment_st' has 5 categories.
 Column 'personal_status' has 4 categories.
 Column 'gurantors' has 3 categories.
 Column 'property_type' has 4 categories.
 Column 'installment_type' has 3 categories.
 Column 'housing_type' has 3 categories.
 Column 'job_type' has 4 categories.
 Column 'telephone' has 2 categories.
 Column 'foreigner' has 2 categories.

Modelling

```
In [30]: # Split data into features (X) and target variable (y)
X = df1_new.drop(columns=['Group_no']) # Features
y = df1_new['Group_no'] # Target variable

In [31]: # Mapping classes from [1, 2, 3] to [0, 1, 2]
class_mapping = {1: 0, 2: 1, 3: 2}
y = y.map(class_mapping)

In [32]: X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=123)

In [33]: # Define classifiers
classifiers = {
    "Random Forest": RandomForestClassifier(n_estimators=100, random_state=0),
    "Decision Tree": DecisionTreeClassifier(random_state=42),
    "Logistic Regression": LogisticRegression(),
    "NB": GaussianNB(),
    "MLP": MLPClassifier(),
    "Support Vector Machine": SVC(kernel='linear', random_state=42),
    "Gradient Boosting": GradientBoostingClassifier(n_estimators=100, learning_rate
}

In [34]: # Define an empty dictionary to store accuracies
accuracies = {}

# Train and evaluate each classifier
for name, clf in classifiers.items():
    print(f"Training {name}...")
    # Train the classifier
    clf.fit(X_train, y_train)
```

```
# Make predictions on the testing set
y_pred = clf.predict(X_test)

# Calculate accuracy
print(" ")

accuracy = accuracy_score(y_test, y_pred)
print(f"          Accuracy for {name}: {accuracy*100:.2f}%")
accuracies[name] = round(accuracy*100,2)

# Print classification report
print(" ")

print(f"Classification Report for {name}:")
print(classification_report(y_test, y_pred))

# Print confusion matrix
print(f"Confusion Matrix for {name}:")
print(confusion_matrix(y_test, y_pred))
print("-----")
```

Training Random Forest...

Accuracy for Random Forest: 67.50%

Classification Report for Random Forest:

	precision	recall	f1-score	support
0	0.71	0.63	0.67	59
1	0.64	0.70	0.67	110
2	0.71	0.68	0.69	71
accuracy			0.68	240
macro avg	0.69	0.67	0.68	240
weighted avg	0.68	0.68	0.68	240

Confusion Matrix for Random Forest:

```
[[37 20  2]
 [15 77 18]
 [ 0 23 48]]
```

Training Decision Tree...

Accuracy for Decision Tree: 49.58%

Classification Report for Decision Tree:

	precision	recall	f1-score	support
0	0.61	0.63	0.62	59
1	0.49	0.45	0.47	110
2	0.42	0.45	0.44	71
accuracy			0.50	240
macro avg	0.50	0.51	0.51	240
weighted avg	0.50	0.50	0.50	240

Confusion Matrix for Decision Tree:

```
[[37 14  8]
 [24 50 36]
 [ 0 39 32]]
```

Training Logistic Regression...

Accuracy for Logistic Regression: 63.75%

Classification Report for Logistic Regression:

	precision	recall	f1-score	support
0	0.71	0.68	0.70	59
1	0.63	0.60	0.62	110
2	0.59	0.66	0.62	71
accuracy			0.64	240
macro avg	0.65	0.65	0.64	240
weighted avg	0.64	0.64	0.64	240

Confusion Matrix for Logistic Regression:


```
[[40 14  5]
 [16 66 28]
 [ 0 24 47]]
```

Training NB...

Accuracy for NB: 60.83%

Classification Report for NB:

	precision	recall	f1-score	support
0	0.58	0.71	0.64	59
1	0.59	0.60	0.60	110
2	0.68	0.54	0.60	71
accuracy			0.61	240
macro avg	0.62	0.62	0.61	240
weighted avg	0.61	0.61	0.61	240

Confusion Matrix for NB:

```
[[42 15  2]
 [28 66 16]
 [ 3 30 38]]
```

Training MLP...

Accuracy for MLP: 57.08%

Classification Report for MLP:

	precision	recall	f1-score	support
0	0.55	0.78	0.65	59
1	0.63	0.35	0.45	110
2	0.55	0.75	0.63	71
accuracy			0.57	240
macro avg	0.58	0.62	0.58	240
weighted avg	0.59	0.57	0.55	240

Confusion Matrix for MLP:

```
[[46  7  6]
 [34 38 38]
 [ 3 15 53]]
```

Training Support Vector Machine...

Accuracy for Support Vector Machine: 65.00%

Classification Report for Support Vector Machine:

	precision	recall	f1-score	support
0	0.68	0.61	0.64	59
1	0.62	0.67	0.65	110
2	0.68	0.65	0.66	71
accuracy			0.65	240

macro avg	0.66	0.64	0.65	240
weighted avg	0.65	0.65	0.65	240

Confusion Matrix for Support Vector Machine:

```
[[36 20  3]
 [17 74 19]
 [ 0 25 46]]
```

Training Gradient Boosting...

Accuracy for Gradient Boosting: 68.75%

Classification Report for Gradient Boosting:

	precision	recall	f1-score	support
0	0.66	0.69	0.68	59
1	0.69	0.65	0.67	110
2	0.71	0.73	0.72	71
accuracy			0.69	240
macro avg	0.69	0.69	0.69	240
weighted avg	0.69	0.69	0.69	240

Confusion Matrix for Gradient Boosting:

```
[[41 14  4]
 [21 72 17]
 [ 0 19 52]]
```

In [35]: accuracies

Out[35]: {'Random Forest': 67.5,
'Decision Tree': 49.58,
'Logistic Regression': 63.75,
'NB': 60.83,
'MLP': 57.08,
'Support Vector Machine': 65.0,
'Gradient Boosting': 68.75}

```
In [36]: # Define a color palette for the markers
color_palette = ['rgb(31, 119, 180)', 'rgb(255, 127, 14)', 'rgb(44, 160, 44)', 'rgb(255, 102, 102)', 'rgb(153, 153, 153)']

# Create a trace for accuracy values with marker colors from the palette
trace = go.Bar(
    x=list(accuracies.keys()),
    y=list(accuracies.values()),
    marker_color=color_palette,
    name='Accuracy'
)

# Create the layout for the graph
layout = go.Layout(
    title='Model Accuracies',
    xaxis_title='Classifier',
    yaxis_title='Accuracy (%)',
    hovermode='closest',
```

```

    plot_bgcolor='rgba(0,0,0,0)'
)

# Create the figure
fig = go.Figure(data=[trace], layout=layout)

# Show the interactive graph
fig.show()

```

Model Accuracies



Selecting our best model

```

In [37]: # Hyperparameter training for Random forest Classifier

# Define parameter grid for hyperparameter tuning
param_grid = {
    'n_estimators': [100, 200, 300],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}

# Initialize and tune Random Forest Classifier
rfc = RandomForestClassifier(random_state=123)

```

```

grid_search = GridSearchCV(rfc, param_grid, cv=5, scoring='accuracy')
grid_search.fit(X_train, y_train)

# Get the best parameters
best_params = grid_search.best_params_

# Train RFC with best parameters
rfc_best = RandomForestClassifier(**best_params)
rfc_best.fit(X_train, y_train)

```

Out[37]:

```

RandomForestClassifier
RandomForestClassifier(min_samples_leaf=4, min_samples_split=10,
                        n_estimators=200)

```

In [38]:

```

# Initialize and train Gradient Boosting Classifier
gbc = GradientBoostingClassifier(n_estimators=100, learning_rate=1.0, max_depth=1,
gbc.fit(X_train, y_train)

# Predict on the test set
y_pred_gb = gbc.predict(X_test)

# Evaluate model performance
accuracy_gb = accuracy_score(y_test, y_pred_gb)
print("Accuracy for Gradient Boosting Classifier:", round(accuracy_gb * 100, 2))

# Classification report
print("Classification Report for Gradient Boosting Classifier:")
print(classification_report(y_test, y_pred_gb))

```

```

Accuracy for Gradient Boosting Classifier: 68.75
Classification Report for Gradient Boosting Classifier:

```

	precision	recall	f1-score	support
0	0.66	0.69	0.68	59
1	0.69	0.65	0.67	110
2	0.71	0.73	0.72	71
accuracy			0.69	240
macro avg	0.69	0.69	0.69	240
weighted avg	0.69	0.69	0.69	240

Even After hyperparameter tuning Gradient boosting works best for us

In [39]:

```
y_pred_gb
```

```
Out[39]: array([0, 1, 1, 0, 2, 1, 2, 1, 2, 0, 0, 0, 1, 2, 1, 0, 2, 2, 0, 0, 1, 2,
                1, 2, 0, 0, 2, 2, 0, 1, 1, 1, 2, 0, 2, 1, 2, 1, 1, 2, 0, 2, 1, 2,
                0, 0, 1, 2, 2, 1, 2, 2, 1, 0, 1, 0, 2, 2, 2, 1, 1, 1, 0, 0, 0, 1,
                2, 0, 2, 2, 2, 0, 0, 1, 2, 0, 2, 1, 1, 1, 1, 2, 0, 2, 0, 1, 1, 0,
                1, 0, 1, 1, 0, 0, 2, 2, 2, 0, 1, 0, 1, 0, 1, 0, 0, 1, 1, 1, 1, 1,
                1, 0, 2, 2, 1, 1, 2, 1, 1, 1, 0, 1, 0, 1, 1, 2, 0, 2, 1, 1, 2, 1,
                1, 2, 2, 2, 1, 1, 1, 1, 1, 1, 2, 0, 1, 0, 2, 1, 1, 0, 0, 1, 1, 0,
                1, 0, 0, 1, 1, 1, 2, 1, 1, 0, 0, 2, 0, 0, 1, 1, 1, 0, 2, 1, 1, 1,
                1, 0, 2, 2, 0, 0, 0, 2, 0, 2, 2, 2, 1, 2, 2, 1, 2, 1, 1, 1, 0,
                1, 1, 2, 2, 1, 2, 1, 1, 0, 1, 2, 0, 1, 1, 2, 2, 1, 2, 2, 1, 1, 0,
                1, 2, 2, 1, 1, 1, 1, 2, 1, 0, 2, 1, 1, 2, 2, 2, 1, 0, 1],
                dtype=int64)
```

Save Model

```
In [40]: import pickle

# Save the model to a file
with open('Credit Classification Problem Model C&T Bank.pkl', 'wb') as file:
    pickle.dump(GradientBoostingClassifier, file)

# Load the model from the file
with open('Credit Classification Problem Model C&T Bank.pkl', 'rb') as file:
    loaded_model = pickle.load(file)
```

Submission file

```
In [47]: print("Training Data Columns:", X_train.columns)
print("Prediction Data Columns:", df2.columns)

Training Data Columns: Index(['sno', 'duration_month', 'poi', 'resident_since', 'age',
                             'credits_no',
                             'liables', 'acc_info', 'credit_history', 'purpose', 'savings_acc',
                             'employment_st', 'personal_status', 'gurantors', 'property_type',
                             'installment_type', 'housing_type', 'job_type', 'telephone',
                             'foreigner'],
                             dtype='object')
Prediction Data Columns: Index(['sno', 'acc_info', 'duration_month', 'credit_history',
                                'purpose',
                                'savings_acc', 'employment_st', 'poi', 'personal_status', 'gurantors',
                                'resident_since', 'property_type', 'age', 'installment_type',
                                'housing_type', 'credits_no', 'job_type', 'liables', 'telephone',
                                'foreigner'],
                                dtype='object')
```

```
In [48]: # Reorder columns in the prediction data to match training data
df2 = df2[X_train.columns]
```

```
In [49]: submission_pred=gbc.predict(df2)
```

```
In [50]: submission_pred
```

```
Out[50]: array([1, 1, 0, 2, 2, 0, 2, 2, 0, 2, 2, 2, 0, 0, 0, 0, 1, 1, 0, 1, 2, 1,
1, 2, 1, 2, 0, 1, 0, 0, 1, 2, 0, 0, 2, 2, 2, 1, 0, 1, 1, 1, 2, 1,
0, 1, 1, 2, 2, 2, 2, 0, 2, 1, 0, 2, 2, 1, 2, 1, 1, 1, 1, 1, 1, 1,
1, 1, 0, 2, 0, 2, 1, 2, 1, 2, 1, 1, 2, 0, 0, 0, 0, 2, 1, 2, 1, 0,
0, 0, 0, 2, 2, 0, 1, 0, 1, 1, 1, 1, 1, 1, 0, 2, 1, 1, 1, 0, 0, 2,
1, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 0, 2, 1, 1, 2, 1, 0, 1, 2, 1, 2,
2, 2, 2, 0, 2, 2, 0, 0, 0, 2, 1, 1, 1, 0, 0, 2, 1, 0, 2, 0, 1, 0,
2, 0, 0, 2, 1, 1, 2, 1, 1, 1, 2, 1, 1, 1, 0, 1, 2, 1, 1, 0, 1, 1,
2, 2, 1, 2, 0, 0, 0, 0, 1, 2, 0, 1, 0, 1, 1, 2, 1, 0, 2, 1, 0, 2,
0, 0], dtype=int64)
```

```
In [51]: copy.head()
```

Out[51]:

	sno	acc_info	duration_month	credit_history	purpose	savings_acc	employment_st	po
0	1	A14	24	A34	A46	A61	A75	4.0
1	2	A12	18	A34	A43	A61	A75	3.0
2	3	A11	20	A34	A42	A61	A75	1.0
3	4	A14	12	A34	A43	A65	A75	4.0
4	5	A12	12	A32	A40	A65	A71	1.0

```
In [52]: final=pd.DataFrame({'serial number':copy.sno.values,'Group_no':submission_pred})
```

```
In [53]: final
```

Out[53]:

	serial number	Group_no
0	1	1
1	2	1
2	3	0
3	4	2
4	5	2
...
195	196	1
196	197	0
197	198	2
198	199	0
199	200	0

200 rows × 2 columns

```
In [54]: # Mapping classes from [0, 1, 2] to [1, 2, 3]
class_mapping = {0: 1, 1: 2, 2: 3}
final['Group_no'] = final['Group_no'].map(class_mapping)
```

```
In [55]: final
```

```
Out[55]:
```

	serial number	Group_no
0	1	2
1	2	2
2	3	1
3	4	3
4	5	3
...
195	196	2
196	197	1
197	198	3
198	199	1
199	200	1

200 rows × 2 columns

```
In [56]: final.to_csv('submission.csv', index=False) #by default index is true
```